

Mammographic Mass Data Set

Sergio Castillo, 1513228@rgu.ac.uk

22nd November 2019

Machine Learning and Big Data Analytics (CM3111)
School of Computing Science & Digital Media
Robert Gordon University

Contents

1	Overview	1
2	Data Exploration	1
2.1	The Dataset	1
2.1.1	Dataset Name	1
2.1.2	Dataset Source	1
2.1.3	Dataset Acquirement Date	1
2.2	Problem Statement & Data Exploration	2
2.2.1	Objective	2
2.2.2	Attributes	3
2.2.3	Class Distribution	4
2.2.4	Dataset Summary	5
2.2.5	Attribute Correlation	7
3	Data Pre-processing	10
3.1	Irrelevant Columns	10
3.2	Missing Values	10
3.2.1	Filling missing values	11
3.2.2	Deleting rows	12
3.3	Missing Header	12
3.4	Categorical Values	13
3.5	Data Normalization	14
4	Modelling	14
4.1	Creating Subsets	14
4.2	Model	15
4.3	Test Model	15
4.4	Results	15
5	Improving Performance	16
5.1	Filling Missing Values	16
5.2	Tune Parameters	16
5.3	Metrics	16
5.4	Different Models	16
5.5	Change Partitioning	16
5.6	Cross Validation	16
6	References	17

1 Overview

The purpose of this report is to analyze a given dataset, in order to reach certain conclusions and be able to predict an outcome with the highest accuracy possible, thus showing competence in understanding, identifying, applying and evaluating machine learning algorithms. In order to do that, I will be using a data set focused on the prediction of malicious mammographic masses based on the result of BI-RAD (Breast Imaging-Reporting and Data System) assessments.

2 Data Exploration

2.1 The Dataset

I've chosen this dataset because I am interested in learning how data analysis and machine learning can be used in order to save people's lives. Additionally, the number of attributes and instances seemed adequate for analysis.

2.1.1 Dataset Name

Mammographic Mass

2.1.2 Dataset Source

The dataset was obtained from the UCI Machine Learning Repository[1]. The origin of the dataset is a paper published by M.Elter, R.Schulz-Wendtland and T.Wittenberg, titled **"The prediction of breast cancer biopsy outcomes using two CAD approaches that both emphasize an intelligible decision process"**[2].

2.1.3 Dataset Acquisition Date

The dataset was obtained on the 4th of October 2019

The original paper was published in 2007, and the dataset was donated on the 29th of October, 2007.

2.2 Problem Statement & Data Exploration

This dataset contains information of 961 patients' BI-RAD assessment, having different parameters of the BI-RAD attributes, patient's age and the result of such assessment. Mammography is the most effective method for breast cancer screening available today. However, the low positive predictive value of breast biopsy resulting from mammogram interpretation leads to approximately 70% unnecessary biopsies with benign outcomes. To reduce the high number of unnecessary breast biopsies, several computer-aided diagnosis (CAD) systems have been proposed in the last years. These systems help physicians in their decision to perform a breast biopsy on a suspicious lesion seen in a mammogram or to perform a short term follow-up examination instead.

2.2.1 Objective

This data set can be used to predict the severity (benign or malignant) of a mammographic mass lesion from BI-RADS (Breast Imaging-Reporting and Data System) attributes and the patient's age.

```
# Dataset dimensions
nrow(df)

## [1] 961

ncol(df)

## [1] 6

dim(df)

## [1] 961 6

# Show some rows
df[1:10,]

##      Assessment Age Shape Margin Density Severity
## 1           5  67     3      5         3         1
## 2           4  43     1      1        <NA>        1
## 3           5  58     4      5         3         1
## 4           4  28     1      1         3         0
## 5           5  74     1      5        <NA>        1
## 6           4  65     1    <NA>         3         0
## 7           4  70    <NA>    <NA>         3         0
## 8           5  42     1    <NA>         3         0
## 9           5  57     1      5         3         1
## 10          5  60    <NA>      5         1         1
```

2.2.2 Attributes

The attributes provide information about the patients' age, mammographic mass results and the severity of the mass. There are 6 attributes in total, from which 4 are predictive attributes, 1 is non-predictive and 1 is the goal field.

1. **Assessment:** Results of the BI-RADS, an indication of how well a CAD system performs compared to the radiologists
 - 0: Definitely benign - 6: Highly suggestive of malignancy
 - Ordinal Categorical Variable
 - Non-Predictive Attribute
2. **Age:** Patient's age in years
 - Integer Variable
 - Predictive Attribute
3. **Shape:** Mass shape
 - 1: Round, 2: Oval, 3: Lobular, 4: Irregular
 - Nominal Categorical Variable
 - Predictive Attribute
4. **Margin:** Mass margin
 - 1: Circumscribed, 2: Microlobulated, 3: Obscured, 4: Ill-defined, 5: Spiculated
 - Nominal Categorical Variable
 - Predictive Attribute
5. **Density:** Mass density
 - 1: High, 2: Iso, 3: Low, 4: Fat-containing
 - Ordinal Categorical Variable
 - Predictive Attribute
6. **Severity:** Severity of the mammographic mass
 - 0: Benign, Malignant: 1
 - Binomial Categorical Variable
 - Goal Attribute

```
# Show Column names
```

```
names(df)
```

```
## [1] "Assessment" "Age"          "Shape"        "Margin"       "Density"
## [6] "Severity"
```

2.2.3 Class Distribution

The goal field on the dataset is the severity of the mammographic mass, which can either be benign (i.e Non-Harmful) or malicious.

```
# Finding how many unique fields there are, and the unique values
length(unique(df$Severity))

## [1] 2

unique(df$Severity)

## [1] 1 0
## Levels: 0 1
```

As shown in the bar plot below, the dataset contains 961 entries, for 516 benign and 445 malignant masses.

```
# Class distribution
severityFreq <- table(df$Severity)
barplot(
  severityFreq,
  col = c('steelblue4', 'coral1'),
  main="Severity Frequency",
  names.arg=c("Benign", "Malicious"),
  ylab="Number of patients"
)
```




Figure 1: Class distribution of the severity

2.2.4 Dataset Summary

The summary displays a briefing of every column, showing statistical information about the data in every entry. The summary can either be numerical (Which shows information such as the maximum and minimum values, median and quartiles) or categorical, which shows the frequency of each one of the categorical values.

In this dataset, only the Patient's age is numerical. The rest of the attributes only show the frequency.

Age Summary

The patient Age attribute is a numerical attribute, so the summary shows the minimum, first quartile, median, mean, third quartile and maximum values of the data

```
# Age Summary  
summary(df$Age)
```

Min	1st Qu	Median	Mean	3rd Qu	Max	Missing
18.00	45.00	57.00	55.49	66.00	96.00	5

```
Age <- df$Age  
hist(  
  Age,  
  xlab="Patient Age",  
  col="turquoise3",  
)
```




Figure 2: Histogram of the Patients' Age

Shape Summary

The mass Shape attribute is a categorical attribute, so the summary shows the value frequency

```
# Shape Summary  
summary(df$Shape)
```

Round	Oval	Lobular	Irregular	Missing
224	211	95	400	31

Margin Summary

The mass Margin attribute is a categorical attribute, so the summary shows the value frequency

```
# Margin Summary  
summary(df$Margin)
```

Circumscribed	Microlobulated	Obscured	Ill-defined	Spiculated	Missing
357	24	116	280	136	48

Density Summary

The mass Density attribute is a categorical attribute, so the summary shows the value frequency

```
# Density Summary  
summary(df$Density)
```

High	Iso	Low	Fat-Containing	Missing
16	59	798	12	76

2.2.5 Attribute Correlation

Correlation measures the relationship between two measurements. This dataset only contains one numerical attribute, so there's no chance to create a numerical correlation between the different values. However, a relationship can be established between the age of the patient and the severity of the mammographic mass. This relation is known, but the other attributes also contribute to it.

```
# Finding how attributes compare to goal
plot(
  data = df,
  Age~Severity,
  col = c('steelblue4','coral1'),
  names.arg=c("Benign","Malicious")
)
```




Figure 3: Correlation between the Age and the severity of the Mammographic mass

In the following page, correlations 3 Age and different Mammographic attributes can be found [4], [5], [6]. As mentioned, there is a relation between the age and the type of mass, but it's not enough to reach a conclusion.

```
library(datasets)
data <- transform(df, Shape)
boxplot(Age~Shape, data, xlab="Shape", ylab="Age")
```




Figure 4: Relation between Age and Mass Shape (Round, Oval, Lobular, Irregular)

```
data <- transform(df, Margin)
boxplot(Age~Margin, data, xlab="Margin", ylab="Age")
```




Figure 5: Relation between Age and Mass Margin (Circumscribed, Microlobulated, Obscured, Ill-defined, Spiculated)

```
data <- transform(df, Density)
boxplot(Age~Density, data, xlab="Density", ylab="Age")
```




Figure 6: Relation between Age and the Mass Density (High, Iso, Low, Fat-containing)

3 Data Pre-processing

In order to analyze the data, the attributes must first be adequated for study. This is known as pre-processing

3.1 Irrelevant Columns

Before going any further, it has been previously mentioned that the Assessment column is a non-predictive attribute, since it measures the effectiveness of CAD Systems. Therefore, the column is not relevant for data mining, and it should be dropped.

```
# Backing up the data frame
df_full <- df

# Dropping the irrelevant column
df$Assessment <- NULL

# Showing the data frame columns without the dropped one
names(df)

## [1] "Age"      "Shape"    "Margin"   "Density"  "Severity"
```

3.2 Missing Values

One of the main issues with the dataset is the number of missing values. When taking a look at the raw data, it can be appreciated that the way missing values were handled was by writing a question mark "?" in place of the value.

The data frame will need to have those missing values recognized, so when parsing the file as a data frame, missing values must be specified

```
# Importing the data
dataFrame <- read.table('data/mammographic_masses.data',
                        header=FALSE, sep=",", na.strings="?")

# Showing some rows
df[1:5,]

##   Age Shape Margin Density Severity
## 1  67     3      5        3         1
## 2  43     1      1    <NA>         1
## 3  58     4      5        3         1
## 4  28     1      1        3         0
## 5  74     1      5    <NA>         1
```

Now that missing values have been identified, it is necessary to know which values are missing:

```
sum(is.na(df))  
## [1] 160  
  
sum(is.na(df$Age))  
## [1] 5  
  
sum(is.na(df$Shape))  
## [1] 31  
  
sum(is.na(df$Margin))  
## [1] 48  
  
sum(is.na(df$Density))  
## [1] 76  
  
sum(is.na(df$Severity))  
## [1] 0
```

Age	Shape	Margin	Density	Severity	Total
5	31	48	76	48	162

3.2.1 Filling missing values

With numerical values, a way to deal with missing values is by replacing the blank attribute with the mean of such attribute. Since the age is the only numerical attribute, we can only use this method here.

```
# Defining the function  
fillMissings <- function(x) {  
  replace(  
    x,  
    is.na(x),  
    mean(x[!is.na(x)]))  
}  
  
# Applying the function to the Age column
```

```
df$Age <- fillMissings(df$Age)

# Verifying the number of missing values on the age column
sum(is.na(df$Age))

## [1] 0
```

3.2.2 Deleting rows

Since we cannot fill the missing values, due to the categorical nature of the attributes, we have no other option but to drop the rows that contain them.

```
# Deleting rows with missing factor attributes
df <- df[!is.na(df$Shape),]
df <- df[!is.na(df$Margin),]
df <- df[!is.na(df$Density),]
```

We can now observe that there are no missing values in our data frame

```
sum(is.na(df))

## [1] 0
```

3.3 Missing Header

Another issue with the dataset is the lack of headers. We may know the meaning of every column after reading about the dataset, but we still need a way to identify each one of the columns.

```
# Showing the data frame header
names(dataFrame)

## [1] "V1" "V2" "V3" "V4" "V5" "V6"

# Changing the column names
names(dataFrame) <- c("Assessment", "Age", "Shape", "Margin", "Density", "Severity")

# Showing the column names
names(dataFrame)

## [1] "Assessment" "Age"          "Shape"       "Margin"     "Density"
## [6] "Severity"
```

3.4 Categorical Values

As previously mentioned, most of the values are categorical, rather than numerical. This means that it can't just be measured as a number, but as properties of the entries. After dropping unnecessary columns, the categorical attributes are the following:

1. **Shape:** Shape of the mammographic mass. (Nominal)
2. **Margin:** Margin of the mammographic mass. (Nominal)
3. **Density:** Density of the mammographic mass. (Ordinal)
4. **Severity:** Severity of the mammographic mass. (Binominal)

Fortunately enough, the categorical attributes already come as numerical values. If this was not the case and they were represented as, for instance, strings, such strings would need to be converted into numerical values.

Now, the attributes will be converted into categorical values.

```
# Turning the attributes into categorical
df$Shape      <- as.factor(df$Shape)
df$Margin     <- as.factor(df$Margin)
df$Density    <- as.factor(df$Density)
df$Severity   <- as.factor(df$Severity)

# Checking whether the fields are factors
is.factor(df$Age)

## [1] FALSE

is.factor(df$Shape)

## [1] TRUE

is.factor(df$Margin)

## [1] TRUE

is.factor(df$Density)

## [1] TRUE

is.factor(df$Severity)

## [1] TRUE
```

3.5 Data Normalization

Since there are no noticeable irregularities in the dataset, such as outliers or skewed values, there is no need to normalize the dataset.

4 Modelling

Now that the dataset has been explored and there is a better idea of what needs to be predicted, we can start building and testing models.

4.1 Creating Subsets

Before creating predictive models, the dataset will be split into a training and a testing set, in order to avoid future problems with overfitting or underfitting.

```
# Downloading necessary libraries
```

```
library(ISLR)
```

```
library(dplyr)
```

```
# Getting the training dataset
```

```
train<-sample_frac( df , 0.7)
```

```
sid<-as.numeric(rownames(train))
```

```
nrow(train)
```

```
## [1] 585
```

```
# Getting the test dataset
```

```
test <- df[-sid,]
```

```
nrow(test)
```

```
## [1] 251
```

```
if ( nrow(train) + nrow(test) == nrow(df) ){
```

```
  print("The dataset has been split properly")
```

```
} else {
```

```
  print("The split of the dataset was not correct")
```

```
}
```

```
## [1] "The dataset has been split properly"
```


4.2 Model

4.3 Test Model

4.4 Results

5 Improving Performance

5.1 Filling Missing Values

Earlier in this coursework, I made the mistake to assume all categorical values were supposed to be deleted, without trying to handle them. However, I realized that there are several algorithms that deal with missing categorical values. So I decided to try them out.

5.2 Tune Parameters

5.3 Metrics

5.4 Different Models

5.5 Change Partitioning

5.6 Cross Validation

6 References

- [1] Matthias Elter. *Mammographic Mass Data Set*. 2007. URL: <https://archive.ics.uci.edu/ml/datasets/MammographicMass>.
- [2] *The prediction of breast cancer biopsy outcomes using two CAD approaches that both emphasize an intelligible decision process*. 2007. URL: https://www.researchgate.net/publication/5776853_The_prediction_of_breast_cancer_biopsy_outcomes_using_two_CAD_approaches_that_both_emphasize_an_intelligible_decision_process.